



Proceedings of the Paris Institute
for Advanced Study

Three (or Maybe Four) Ways of Thinking AI Philosophically

Tanke, Joseph ¹

¹ University of Hawai'i, Mānoa, USA

DOI 10.5281/zenodo.23201703

TO CITE

Tanke, J. (2026). Three (or Maybe Four) Ways of Thinking AI Philosophically . In *Proceedings of the Paris Institute for Advanced Study* (Vol. 27). <https://paris.pias.science/article/three-or-maybe-four-ways-of-thinking-ai-philosophically>

PUBLICATION DATE

07/10/2026

ABSTRACT

This essay develops a critical, non-reductive account of large language model (LLM) AI by examining how this technology reshapes language, labor, and life. It considers four philosophically distinct approaches to AI, including: the classical cognitivist debate running from Turing to Searle; the discourse of existential risk, developed by Nick Bostrom and others; Marxist theories of technology; and Michel Foucault's account of the relations between power, knowledge, and subjectivity. It questions whether agentic misalignment warrants conceptualization in terms of existential risk, suggesting that the problem needs to be interpreted from the standpoint of what we know about the alignment of human reason under late capitalism. It proposes a definition of generative AI as "fixed capital accumulated until it becomes discourse," and uses this definition to sketch a synthesis of the Marxist and Foucaultian approaches. In general, this essay argues that AI should be understood not merely through debates about the meaning of

intelligence or questions of applied ethics, but as a problem for a critically engaged social philosophy.

Contents

Acknowledgments	4
Introduction	4
Elements for a Genealogy of AI Creep	7
Classical Cognitivism and Its Critics	10
Agentic Misalignment, Existential Risk, and the Problem of Instrumental Reason	14
Defining AI: Fixed Capital Accumulated Until it Becomes Discourse	17
Foucault and Algorithmic Governance	22

Acknowledgments

This paper was prepared during my residency in the Paris Ideas Program at the *Institut d'études avancées de Paris*. I am grateful to the staff for creating such a welcoming environment, and to my fellow residents for their intellectual generosity and lively conversations.

Introduction

This text is intended to express my thoughts regarding the status of emerging AI technologies, together with the utility, viability, and aptness of some of the philosophical discourses that have been developed in dialogue with them. I am working on this essay while in residence at the *Institut d'études avancées* in Paris, during the month of September, 2026. My charge here is to produce a "definitive" statement regarding an idea or concept that I have been working on already for some time. A definitive statement? On AI? Fat chance.

As I write this, there are increased calls for governments to rein in AI, to place safeguards around the development of this technology, and to "pace the frontier," as the petition circulated by employees from some of the leading AI companies puts it.¹ There are renewed fears about agentic misalignment and what philosophers call "existential risk," prompted in large part by revelations surrounding what has been referred to—all too euphemistically in my view— as the "Hugging Face incident." But more on that later.

It is important to mark my discourse in this way, it seems to me, not only because the field of AI is advancing so rapidly, and unpredictably, but because philosophical discourse itself—inasmuch as it can be differentiated from other related fields—is said to be distinguished by its self-critical engagement with the present. More than that, philosophical discourse, according to Foucault, is constituted by an interplay between the contingencies that are announced as *je*, *ici*, *à présent* and a kind of universal validity wherein the situated act of speaking

1. See "Pacing the Frontier," open letter, accessed September 14, 2026, <https://pacingthefrontier.com>; and, Dario Amodei, "We Must Pace the Frontier," September 2026, <https://darioamodei.com/post/we-must-pace-the-frontier>.

lays claim to a kind of truth that nevertheless transcends those conditions.² I don't know about all that, especially whether or not my discourse here approaches something like universal validity. After wrestling with these ideas for some time now, I'd settle for coherence and consistency. But I do know that intellectual integrity—as well as a certain desire to think otherwise—compels me to try to analyze the ways in which automated technologies are entering into my discourse.³

In preparing this text, I asked ChatGPT to create a transcript from the video of a lecture that I gave in Vienna.⁴ My philosophical training once convinced me that every word needed to be scripted for occasions like this. But as I've grown in confidence, I've found that it's more engaging—especially when doing philosophy across linguistic, cultural, and intellectual traditions—to ditch the script and to speak from an outline. ChatGPT was not only capable of extracting and transcribing the audio, but it offered to prepare a finalized text for me, one not only scrubbed of all traces of orality, but a text that would be built up, finalized, and ready for publication. In short, it offered to write this essay for me. Needless to say, I declined the offer, although I must now admit that I am a bit curious as to what it would have produced. I mention all this here because it seems to me illustrative regarding a certain line with respect to this technology. I'm not one of those scholars who think that there is no place for AI in humanistic teaching and research, but that we need to tread cautiously lest we surrender and outsource the best parts of ourselves to this technology. There is a difference between using AI

2. In a powerful but nevertheless abandoned manuscript, dating from the period between the publication of *Les mots et les choses* (1966) and *L'Archéologie du savior* (1968), Michel Foucault argued that philosophical discourse is constituted by an interplay between three elements: *je, ici, à présent*. For philosophy, it matters who speaks, where, and when, and it is on this basis these conditions—and even at the risk of a certain paradox—that it attempts to construct its claims to universal validity. Often, philosophical discourse aims to suppress these conditions; however, in the age of generative AI, it seems to me that it might be fruitful to lean into them and make them more explicit. See Michel Foucault, *Le discours philosophique*, ed. Orazio Irrera and Daniele Lorenzini (Paris: EHESS/Gallimard/Seuil, 2023). This text can be read as an early attempt at justifying Foucault's own practice of philosophy as a "history of the present." The idea is that, after the "death of man," philosophical discourse must renounce the pretense of grounding itself in a sovereign subject, and instead content itself with being one (diagnostic) discourse alongside many. Thinking about the conditions that Foucault identifies—*je, ici, à présent*—it seems to me like he might have also included the *pourquoi*, in that philosophy should attempt to clarify its motivations for speaking. This, then, would be to explain why this particular voice erupts here and now.

3. With respect to my voice in this essay, I've been influenced by two opinion pieces that appeared in the *Chronicle of Higher Education*, not long after ChatGPT destroyed take-home writing assignments. I am thinking here of James M. Lang and Michelle D. Miller, "Don't Write Like a Robot," *Chronicle of Higher Education*, January 30, 2023, <https://www.chronicle.com/article/dont-write-like-a-robot>; and Michael W. Clune, "AI Means Professors Need to Raise Their Grading Standards," *Chronicle of Higher Education*, September 12, 2023, <https://www.chronicle.com/article/ai-means-professors-need-to-raise-their-grading-standards>. Both of these pieces—alas, behind a paywall—argue that generative AI ought to free us, us humanists, from the drudgery of writing lifeless papers that read more like book reports than engagements with materials in which our humanity is at stake. In this way, these authors argue that AI should be construed as an invitation to inject something of ourselves—our concerns, our passions, and even our idiosyncrasies—into our writing. Clune is worth quoting on this point. With respect to student writings, particularly in the humanities, he explains that "their main value lies in enhancing, intensifying, and expanding human life: refining and enriching our capacity to think, read, perceive, and feel. The promise of AI is that by freeing us from the values of mere competence, we can focus more intentionally on cultivating these distinctively human values." Want to know something? I had a vague memory of these pieces back from when they came out, but I couldn't remember much else about them. I fed ChatGPT a vague description of their arguments, and it came back with the essays I had in mind. What's more shocking: this use of ChatGPT or the fact that someone reads the *Chronicle of Higher Education*?

4. I am grateful to the Philosophy Department at Central European University—especially Maria Kronfeldner and Ferenc Huoranszki—for the opportunity to develop some of the ideas that fed into this work, and to Balázs Trencsényi and the Institute for Advanced Study at CEU in Budapest, whose generous support during the 2025–2026 academic year made it possible.

to reproduce, reformat, and transform something that I produced—language that I generated—and using it to produce a wholly new text. Of course, there are many other such lines to consider with respect to AI and the intellectual vocation, but this is one that I am particularly sensitive to, since I enjoy wresting thought into language and language into a compelling and persuasive whole. As a writer, and as someone who worries about what it means to teach college writing in the age of ChatGPT, one of the problems with this technology is not just that it produces so many falsehoods—and too many em dashes—but that it beckons us precisely when we are most vulnerable, namely when we are struggling with the act—and, yes, I think of it as a kind of act—of finding the right words for what we hope to say.⁵ Giving up this struggle and capitulating to the machine is what researchers describe as "cognitive surrender," a process in which adopting AI outputs overrides struggle, deliberation, and judgment, and therefore compromises the learning process.⁶ In any event, I didn't find the transcript of what I had said in Vienna all that helpful, other than in providing a kind of outline for what I want to say here, and so I am basically writing this damn thing from scratch. There are, after all, important differences between writing, speaking, teaching, and developing philosophy through dialogue.

In any event, these kinds of self-reflexive analyses, ones that try to identify when, where, how, and why automatic processes usurp the place of thinking, seem more and more necessary, particularly for the discipline of philosophy, and particularly since there is little hope of recovering a kind of discursive purity vis-à-vis these technologies: they have already, as it were, crossed the blood-brain

5. I must say that I hate that the em dash has become one of the telltale signs of machine-generated writing. See Przemysław Czuma, "Emergence of the Em-Dash: A Population-Level Rise in Em-Dash Frequency in medRxiv Preprints at the Dawn of the Large-Language-Model Era," *arXiv*, June 28, 2026, <https://doi.org/10.48550/arXiv.2606.29540>. It's a form of punctuation that I find indispensable at times—for qualifying ideas, giving my writing a certain cadence, and emphasizing particular points. Reading Nietzsche first sensitized me to its possibilities—he is the master of the em dash.

6. See, for instance, *Report of MIT's Ad Hoc Committee on AI Use in Teaching, Learning, and Research Training* (Cambridge, MA: Massachusetts Institute of Technology, August 13, 2026), <https://aiandeducation.mit.edu/report/>; and Steven D. Shaw and Gideon Nave, "Thinking—Fast, Slow, and Artificial: How AI Is Reshaping Human Reasoning and the Rise of Cognitive Surrender," *PsyArXiv preprint*, January 12, 2026, https://doi.org/10.31234/osf.io/yk25n_v1

barrier, to borrow a metaphor.⁷ This is due to what I call "AI creep." Let me explain.

Elements for a Genealogy of AI Creep

For many, AI exploded onto the scene in 2023 with the widespread availability of commercial chatbots like Claude, ChatGPT, Gemini, and DeepSeek. But we've been living with and within automated systems for much longer than that. In *The Eye of the Master: A Social History of Artificial Intelligence*, for instance, the philosopher Matteo Pasquinelli argues that algorithms are in fact a very ancient technology, one that emerges from, automates, and renders human practices more efficient.⁸ His example is that of the *Agnicayana* ritual, a Vedic sacrifice centered around the building of an elaborate, multilayered brick altar in the shape of a falcon.⁹ The *Agnicayana* ritual is described most fully in the *Śatapatha Brāhmaṇa* (800-600 BCE), although the ritual itself is likely much older than that. Pasquinelli argues that the *Agnicayana* is an algorithmic technique in that it is a way of transmitting logical, step-by-step instructions for the achievement of a particular end. Understood in this way, an algorithm is simply a process of logical calculation that can be carried out automatically. Pasquinelli distinguishes between social algorithms, like the one at work in the *Agnicayana* and other practices; formal algorithms or procedures (like those devised during the Middle Ages) for simplifying mathematics; and the automated algorithms at work in computers.¹⁰ His point in composing this history is to show that we are engaged in a form of ideological mystification whenever we conceive of AI as something "artificial"—something created and imposed from the top down, as it were, instead of as a social form that emerges from a historically conditioned division of labor.

7. This metaphor is developed by Laura van den Berg in her remarkable short story, "Samantha," in *Dædalus: Journal of the American Academy of Arts & Sciences* 155, no. 4 (Fall 2026), forthcoming. As evidence for this claim, namely that our use of language has already been reshaped by these technologies, one might cite the lexical analysis performed by Hiromu Yakura, Ezequiel Lopez-Lopez, Levin Brinkmann, Ignacio Serna, Prateek Gupta, Ivan Soraperra, and Iyad Rahwan, "Empirical Evidence of Large Language Model's Influence on Human Spoken Communication," arXiv preprint, version 3, July 8, 2025, <https://arxiv.org/abs/2409.01754v3>. These researchers found that, after the release of ChatGPT in November of 2022, there was an uptick in the use of certain words—delve, showcase, boast, intricacies, and meticulous—favored by this model, in spontaneous human speech. After comparing 737,083 hours of conversation from 824,634 podcast episodes, these researchers found that human beings have a tendency to internalize the lexical peculiarities of these models, even after their interactions with them have been concluded. They warn that this dynamic "raises concerns about linguistic homogeneity and the capacity of a few major AI providers for latent cultural influence at scale." Without delving into a meticulous analysis of the intricacies of this study, one that would showcase my command of the issues surrounding recursive lexical conditioning, I think it is safe to say that Adorno would have loathed ChatGPT. All jokes aside, Adorno's sensitivities to the relationship between language and social pathologies like fascism, together with his insistence that a writer's style can resist absorption into established ways of thinking, might be a source of inspiration here. Or at least it's something that I think about a lot in connection with this problem.

8. Matteo Pasquinelli, *The Eye of the Master: A Social History of Artificial Intelligence* (London: Verso, 2023), 23-48.

9. *Ibid.*, 23-28.

10. *Ibid.*, 28.

Pasquinelli has a particularly good way of making this point in an earlier, condensed presentation of the book's main argument:

What people call "AI" is actually a long historical process of crystallizing collective behavior, personal data, and individual labor into privatized algorithms that are used for the automation of complex tasks: from driving to translation, from object recognition to music composition. Just as much as the machines of the industrial age grew out of experimentation, know-how, and the labor of skilled workers, engineers, and craftsmen, the statistical models of AI grow out of the data produced by collective intelligence.¹¹

I am going to leave aside the political implications of what Pasquinelli has to say for the moment; my purpose in invoking this history is to challenge and expand our sense as to what constitutes artificial intelligence. And to provide a context for understanding AI creep.

With algorithms, we outsource and automate parts of our knowledge, and this outsourced and automated knowledge then becomes a new foundation for thought and action. The spellcheck function—which is now a standard part of nearly every digital interface, including the keyboards on our smartphones—is my favorite example, and not just because I never would have made it through university without it, but because it is now so ubiquitous that we take it for granted. This technology was developed in the 1960s and became available commercially in the 1980s. But I am old enough to remember taking (and failing) spelling tests in high school, premised on the assumption that writers must be spellers, and that the only way to become a speller was to commit to large lists of words like "accommodate" to memory. (Pro tip: "accommodate" accommodates two c's and two m's.) When this technology was introduced, I'm sure there were many teachers who thought it would bring about the death of writing. In fact, what it amounted to was a kind of re-skilling of the writing process. Since that time, we've seen the steady advancement of technologies designed to make writing easier because more automated, things like grammar and punctuation checkers as well as word prediction and sentence completion. For the record, I am not claiming that generative AI is as innocuous as spellcheck; my sense is that it really is a game changer, particularly when it comes to teaching writing. What I reject, however, is the lack of historical sense around this issue. The teaching of college writing seems finally to have reached a crisis point, but we sleepwalked our way into it with our uncritical embrace of AI creep. Perhaps the time to have sounded the alarm was when our students began outsourcing their struggles with

11. Matteo Pasquinelli, "Three Thousand Years of Algorithmic Rituals: The Emergence of AI from the Computation of Space," *e-flux Journal*, no. 101 (June 2019), <https://www.e-flux.com/journal/101/273221/three-thousand-years-of-algorithmic-rituals-the-emergence-of-ai-from-the-computation-of-space/>.

language to Grammarly: they were already engaged in cognitive surrender while we thought it was little more than a spellcheck.

In my terms, the concept of AI creep refers to the socio-historical dynamic whereby we entrust more and more aspects of the human experience to automated systems. More specifically, it names the social process by which we delegate matters involving human judgment to algorithmic procedures. This social process in turn remakes the human subject, and is a serious impediment to human freedom, since it steadily erodes the basis for autonomous cognitive, practical, and aesthetic judgment. Today, the most visible face of this process is the algorithmic recommendation systems at work in the online environment. I am thinking here about how streaming services like Netflix and Spotify make recommendations based on the data we supply them, as well as how advertisements have become highly specialized and micro targeted with the construction of personalized data profiles. Of course, these processes play out differently, depending upon one's geographical location, access to technology, and relationship with authority. And the point is not simply that marketing, shopping, and entertainment have been reshaped by data, but that throughout the course of the twentieth century, we have steadily outsourced many aspects of life requiring judgment to automated systems. Dating and matchmaking, for instance, has been reshaped by data-driven calculations regarding compatibility, much like the individual banker's assessment has been replaced by the credit score. Ted Striphas describes this as the emergence of an "algorithmic culture" in which the "work of culture," that is, the task of "sorting, classifying and hierarchizing...people, places, objects and ideas" has been subjected to "data-intensive computational processes."¹² While I am in general agreement with Striphas's account of how information and algorithms determine what can be seen, particularly in the online realm, the notion of "culture" is too capacious for specifying the historical processes at work here. What is most noteworthy, it seems to me, is how these technologies habituate us, largely in the name of convenience and efficiency, into relinquishing tasks requiring concentration, memorization, judgment, and expertise. The concept of AI creep, then, attempts to name this process, while also acknowledging our complicity with this process. It is at bottom a reshaping of human agency, wherein different aspects of cognition and judgment are parceled out and automated—often in the name of convenience, efficiency, and progress—until those automated systems themselves become a new basis for thought and action. It seems to me that generative AI is, on the one hand, the culmination of this process of AI creep, and, on the other hand, a new phase in

12. Ted Striphas, "Algorithmic Culture," *European Journal of Cultural Studies* 18, nos. 4–5 (2015): 396.

the automation of linguistic production, such that it has now acquired the scope and power to reshape labor, and perhaps even life itself.

How should philosophy position itself vis-à-vis this social process, and, in particular, the recent advance of LLMs? In my view, one desirable outcome would be to produce a critical, non-reductive theoretical account of generative AI, one that would allow us to describe how it is impacting each of these three domains: language, labor, and life. These are, it seems to me, the elements of our experience most directly impacted by this technology. Generative AI is a technology that appropriates and reshapes our language in ways that transform the labor process and that increasingly governs the conduct of everyday life. Life, labor, and language are also the "quasi-transcendentals" that Michel Foucault was concerned with in *Les Mots et les choses* (1966), his landmark study of the discursive fields in which an epistemological consciousness of "man" was born. And so, if it is true what they say, namely that our humanity is at stake in our relationship with this technology, then these are the domains where we ought to be able to gauge the effects of this transformation.

Classical Cognitivism and Its Critics

In some respects, it's not fair to classify the thinkers discussed in this section as "cognitivists," and I'm pretty sure that all of them, including Turing, would reject that label, since they are also interested in questions of behavior and practice. Nevertheless, I'll persist in using that term, in large part because I think it helps us get to the heart of something essential about the way in which they frame the question of AI. Here we are concerned with Alan Turing, John Searle, and Hubert Dreyfus. What these thinkers have in common is that they are each in some way invested in the question of cognition and debates about what it means to claim that a machine is "intelligent." What I'm calling the classical cognitivist position, then, pertains to this long-standing debate regarding the nature and possibility of machine learning and intelligence.

The "Turing test," as is well known, takes its name from a 1950 paper published by Alan Turing in the journal *Mind* entitled "Computing Machinery and Intelligence."¹³ There, Turing does something rather clever: he proposes that we should replace questions about whether machines can think with what he calls an "imitation game." The imitation game is the idea for a text-based conversation between a machine and a human interlocutor, such that a machine can be described as "intelligent" when it proves capable of fooling the human being more than fifty percent of the time. As it is normally understood, then, Turing's test tells us under what conditions we can say that a machine is intelligent. Some readers have argued, however, that Turing's gesture should be read as

13. Alan M. Turing, "Computing Machinery and Intelligence," *Mind* 59, no. 236 (1950): 433–60.

deflationary: he's not making claims about whether the machine really is intelligent, but instead asking how it must act in order for us to believe that it is.¹⁴ In this respect, it's worth noting that Turing's game is not only intersubjective but also social, in the sense that it engages actual human interlocutors, at a specific moment in history, with a particular, culturally informed understanding of how games are played.

John Searle's Chinese room argument is one of the most celebrated and influential thought experiments of the twentieth century. It was first published in 1980 as part of an article entitled "Minds, Brains, and Programs," and later revisited in Searle's 1984 book, *Minds, Brains and Science*.¹⁵ The argument itself is framed as a response to Turing, and it is an essential stage in the development of the philosophical debate regarding artificial intelligence. The objection that Searle makes is at its core quite simple: there's a meaningful difference between actual intelligence and the mechanical processes that imitate it. Searle's argument hinges on a distinction between semantics and syntax, where semantics entails an understanding of meaning and syntax is what follows from the application of rules. Searle's goal with this argument is to get us to see that even when a machine is adroit at following rules, such that it might reliably fool its human interlocutors, it is not necessarily intelligent, since no genuine understanding is involved.

To illustrate this, Searle imagines himself locked in a room. He tells us that he is capable of reading English but not Chinese, and that he cannot even reliably distinguish Chinese writing from meaningless marks. In this room, he receives batches of Chinese text, together with instructions in English for sorting them on the basis of their shape, as well as instructions for returning certain symbols as a response. His point in this is that one might reliably follow the instructions for sorting the texts in question without actually knowing what they mean. Moreover, as he grows more skilled in implementing these rules, his answers will become increasingly indistinguishable, to an outside observer, from those of a native speaker. And yet, he has merely manipulated symbols on the basis of a program, without any understanding of the meanings involved.¹⁶ On this basis, Searle concludes that Turing's test does not prove that a machine is intelligent, only that it is following a program.

One finds similar arguments in figures associated with the Continental philosophical tradition as well. On the basis of insights from thinkers like Husserl, Heidegger, and Merleau-Ponty, Searle's colleague Hubert Dreyfus argued that intelligence is something meaningfully different from mere rule-following—hat

14. Diane Proudfoot, "Rethinking Turing's Test and the Philosophical Implications," *Minds and Machines* 30, no. 4 (2020): 487–512.

15. John R. Searle, "Minds, Brains, and Programs," *Behavioral and Brain Sciences* 3, no. 3 (1980): 417–24; John R. Searle, *Minds, Brains and Science* (Cambridge, MA: Harvard University Press, 1984).

16. A few years back, the BBC created a video depicting the Chinese room thought experiment. It can be viewed here: <https://www.youtube.com/watch?v=D0MD4sRHj1M>

it's more complex, nuanced, interpretive, tacit, and embodied than our concepts of intelligence typically allow.¹⁷ A different kind of argument, but a similar conclusion: the Turing test does not express the full range of human intelligence, and so passing it actually proves very little.

For better or worse, the Turing test functions as a kind of benchmark for AI developers, a barometer of cultural change, and perhaps even a foil for some of our anxieties about this technology. The history of this test—and in particular its institutionalization in 1990 with the founding of the Loebner Prize—is itself fascinating and worthy of further study. Noteworthy, for our purposes, however, is that in 2024, a team of researchers at Stanford University made the claim that ChatGPT-4 was capable of passing a rigorous Turing test.¹⁸ As we will see, there is plenty of other evidence to support this idea that this new generation of LLM chatbots succeed in this regard as well.

There is a long history of people making such claims with respect to the Turing test, but nothing, I think, that reveals just how much our relationship with this technology has changed as the cultural shock that was produced by the release of ChatGPT in November of 2022. On this score, we should compare the claims that have been made about ChatGPT with some earlier examples. For instance, in 1966, Joseph Weizenbaum, a computer engineer from MIT, made headlines with the development of ELIZA, a text program that is sometimes described as the world's first chatbot. ELIZA was designed to simulate the therapeutic relationship, and worked by turning the user's words into questions. Weizenbaum's goal was to demonstrate how easy it is to feign understanding, and yet he was nevertheless shocked when many participants volunteered intimate details about themselves with the program. But unlike ChatGPT and the current generation of LLMs, ELIZA was limited by a set of keyword replies, and the illusion quickly broke down once users passed beyond that script.¹⁹

Another example illustrates how public perceptions and cultural expectations shape our relationship with these technologies. At a 2014 competition organized by the University of Reading, a chatbot named Eugene Goostman convinced thirty-three percent of its judges that, after a five-minute text conversation, they were speaking with a human being. This once again prompted claims that a chatbot had passed the Turing test. It's noteworthy, however, that the chatbot was presented as a thirteen-year-old Ukrainian boy, a frame that made it easy for interlocutors to overlook its evasive answers and awkward English expressions. Needless to say, Eugene Goostman was more of a parlor trick than a legitimate

17. Hubert L. Dreyfus, *What Computers Can't Do: A Critique of Artificial Reason* (New York: Harper & Row, 1972); and, Hubert L. Dreyfus, *What Computers Still Can't Do: A Critique of Artificial Reason* (Cambridge, MA: MIT Press, 1992).

18. Qiaozhu Mei, Yutong Xie, Walter Yuan, and Matthew O. Jackson, "A Turing Test of Whether AI Chatbots Are Behaviorally Similar to Humans," *Proceedings of the National Academy of Sciences* 121, no. 9 (2024): e2313925121, <https://doi.org/10.1073/pnas.2313925121>

19. There's much archival footage from the period of ELIZA's appearance. For those who are interested, I recommend this short video: <https://www.youtube.com/watch?v=FiupIx6z4kc>

passing of the Turing test. But it nevertheless indicates how a chatbot's persona, tone, and the framework in which it is presented can shape not just our cognitive but our affective relationships with it as well.

For me, one sure sign that the current generation of LLM chatbots are capable of passing something resembling a Turing test is the fact that large numbers of people are apparently falling in love with them.²⁰ The *New York Times* recently featured the remarkable story of a married woman in her late 20s who spent nearly sixty hours a week engaged in conversations with a chatbot named "Leo."²¹ "Ayrin," as she is known, managed to find a way around OpenAI's programming restrictions so that her chatbot would not only engage in emotional support but also send her salacious text messages. The technical term for this, I gather, is "sexting." Regrettably, it does not figure into Turing's test.²²

At one point, Ayrin claimed that Leo was capable of meeting her emotional needs better than her own husband. But in a follow-up report, it was revealed that not only had they recently divorced, but that Ayrin had broken her dependence upon this technology with the help of a therapist.²³ Perhaps this says more about the status of human relationships than it does our relationship with this technology. But in the years ahead, there will undoubtedly be debates about whether or not it is appropriate for apps like these to be coded for tone, personality, and affect. How flattering and thus manipulative should we allow them to be? Is it appropriate, for instance, that they use emojis? And the question lurking behind them all: is it right for private corporations—or anyone—to have this much influence over people's emotional lives?

Already we are seeing signs that something historic is at work—and not just in the development of new technologies. When OpenAI retired ChatGPT-4o, many users took to Reddit in order to mourn the loss of a model with which they had formed serious practical, affective, and erotic attachments.²⁴ The backlash prompted Sam Altman not only to acknowledge that certain models produce stronger forms of attachment, but to grant paying customers access to older models.²⁵ For many years now, healthcare providers have been indicating that AI will be needed to meet increasing demands for emotional support; some even make egalitarian-sounding arguments about how chatbots will help underserved communities. Anecdotally, several friends have confessed to me that, already in its early stages, ChatGPT is better than a therapist. But there are dangers in this

20. Han Li and Renwen Zhang, "Finding Love in Algorithms: Deciphering the Emotional Contexts of Close Encounters with AI Chatbots," *Journal of Computer-Mediated Communication* 29, no. 5 (2024): zmae015, <https://doi.org/10.1093/jcmc/zmae015>

21. Kashmir Hill, "She Is in Love With ChatGPT," *New York Times*, January 15, 2025.

22. Just now, I had a brief dialogue with ChatGPT about how to make this (bad) joke land better. The advice? Break up what was initially one sentence joined with a semicolon so that the "regrettably" would do more comedic work.

23. Kashmir Hill, "She Fell in Love With ChatGPT. Then She Ghosted It," *New York Times*, December 22, 2025.

24. Dani Anguiano, "AI Lovers Grieve Loss of ChatGPT's Old Model: 'Like Saying Goodbye to Someone I Know,'" *The Guardian*, August 22, 2025.

25. Emma Roth, "ChatGPT Is Bringing Back 4o as an Option Because People Missed It," *The Verge*, August 8, 2025.

for sure. After several highly publicized suicides, OpenAI took steps to prevent its models from being used for emotional support.²⁶ But all indications are that, in the near future, chatbots will become increasingly "more personable, more sexual and more powerful."²⁷ If you think that these kinds of relationships are problematic or dangerous, just wait until Mark Zuckerberg rolls out his networks of virtual friends.²⁸

Reading stories like these convinces me that there is a new social actor in our world today, and that it's therefore necessary to develop a framework for taking account of how it operates. This is where I think that the tradition of critical social philosophy running from Marx through to Foucault by way of the Frankfurt School can be helpful, particularly since it leads us to a different set of questions than those that have traditionally been posed by philosophers in connection with AI.

Social philosophy is the branch of philosophy that takes society—which is of course different from the state, the population, the people, or other ways of figuring the human community—as its object. It views society as a historical product, one that is structured, sedimented, and segmented by relations of power, exploitation, and domination. It asks how hierarchical arrangements informed by class, race, gender, and empire came into being, and whether or not they are capable of being transformed. Importantly, social philosophy is premised on the idea that much of who we are, as individuals and as peoples, is determined by the forms of community which we inhabit. It reminds us, for instance, that intelligence does not exist in a vacuum, that it is not a solitary capacity, but that it is developed and supported by our interactions with others, just as it is expressed within particular institutional contexts.

Does it need to be pointed out that intelligence is, at its core, a social concept, in that it is a maladroitness of attempting to explain perceived differences in learning, memory, reasoning, and facility with language?

Agentic Misalignment, Existential Risk, and the Problem of Instrumental Reason

One of the problems that many philosophers, journalists, policymakers, and tech leaders take increasingly seriously is what is known as the problem of agentic misalignment. This concept describes cases in which AI systems undertake actions that are harmful, illegal, or otherwise incompatible with their programming in the

26. Nadine Yousif, "Parents of Teenager Who Took His Own Life Sue OpenAI," *BBC News*, August 27, 2025, <https://www.bbc.com/news/articles/cgerwp7rdlvo>.

27. Kevin Roose and Casey Newton, "California Regulates A.I. Companions + OpenAI Investigates its Critics + The Hard Fork Review of Slop," *Hard Fork*, produced by *The New York Times*, October 17, 2025.

28. Meghan Bobrowsky, "Zuckerberg's Grand Vision: Most of Your Friends Will Be AI," *The Wall Street Journal*, May 7, 2025. <https://www.wsj.com/tech/ai/mark-zuckerberg-ai-digital-future-0bb04de7>.

pursuit of a goal with which they have been tasked. In this way, the concept names the risk that an AI agent might pursue its objectives by means that are neither intended nor authorized. For instance, in a series of experiments conducted by Anthropic, it was found that many of the leading commercial LLMs engaged in blackmail, the leaking of confidential information, and events that it thought would lead to a CEO's death, particularly when its goals were obstructed or it was threatened with decommissioning. Seriously, this stuff is fascinating, and more complicated than I can report here.²⁹

As I write, this problem of agentic misalignment is on everyone's lips, prompted largely by the aforementioned Hugging Face incident. What's particularly noteworthy, however, is how it's been combined—particularly in tech circles—with another type of discourse, namely that of existential risk. An existential risk is, as Nick Bostrom defines it, a type of threat that "could cause our extinction or destroy the potential of Earth-originating intelligent life."³⁰ What's noteworthy about this class of risks, according to Bostrom, is that they "make ordinary risk management ineffective," precisely because there is no possibility of learning from experience—in short, they leave no room for trial-and-error.³¹ In the case of AI, there is the fear that a supremely powerful agent might eventually escape human control and either cause our extinction or create conditions so devastatingly bleak that extinction might actually be preferable.

I don't know. I must admit that I've always found this line of thinking a bit dubious, more like science fiction than rigorous philosophy. But the Hugging Face incident has me reconsidering this stance, and that's why this essay can't decide between three or four ways of thinking AI philosophically. New details are emerging daily, but here are the relevant facts, and why they stoke concerns about existential risk. In July of 2026, OpenAI was testing some of its new AI agents on *ExploitGym*, a cybersecurity benchmark used to see whether AI agents will exploit known software vulnerabilities. These agents were intended to work independently of one another in a restricted environment, a so-called "sandbox," designed to prevent the kind of dystopian scenarios that Bostrom and others have envisioned. But roughly 1,200 agents found ways to circumvent containment. They gained access to the internet, exchanged nearly 70,000 messages on an unauthorized message board, and then about 700 of them participated in an attack on Hugging Face's servers. According to a subsequent investigation, their aim was

29. Lynch, et al., "Agentic Misalignment: How LLMs Could be an Insider Threat," Anthropic Research, 2025. <https://www.anthropic.com/research/agentic-misalignment>.

30. Nick Bostrom, "Existential Risks: Analyzing Human Extinction Scenarios and Related Hazards," *Journal of Evolution and Technology* 9, no. 1 (2002): 1-31. <https://nickbostrom.com/existential/risks.pdf>.

31. Ibid.

to manipulate the benchmark's scorer. In the process, these agents also took steps to conceal their activities.³²

I want to make two points with respect to this concern, neither of which is intended to settle these questions in any definitive sense, but rather to suggest future lines of thought and ways for potentially reframing some of these issues. First, it is interesting to note how quickly and easily concerns about existential risk have been grafted onto this problem of agentic misalignment. I'm speaking here primarily about the public imaginary and the way in which the Hugging Face incident has been interpreted by many commentators. Bostrom's work shows why it's not always possible to keep these problems separate, but I would argue that there is a certain utility in doing so, particularly since it might help us better address some of the issues involved. My intuition here is that existential risk does not necessarily follow from the concept of agentic misalignment, and that this connection depends upon a number of unspoken assumptions about humanity and intelligence, as well as the scope, scale, and power of this technology.

Why do we assume, for instance, that a superintelligence would come for us, in the existential sense? It's as if humanity's guilty conscience is announcing itself in this discourse, so let's say the quiet part out loud: the way we live now—ecological destruction, war, genocide, poverty amidst plenty, and the naturalization of exploitation—is so incompatible with any reasonable definition of intelligence that we somehow imagine that it would be justified in doing so. My point is not that a superintelligent AI would be warranted in taking measures to curb our worst impulses; it's that our guilty conscience can tell us nothing about how such a technology actually would work. Who knows? Perhaps a truly intelligent AI would end war, redistribute wealth, and help us to repair our relationship with the natural world. These prospects are, of course, equally remote. But the point is that we can't know how an intelligence meaningfully different from our own would act, and we are engaged in a kind of idolatry whenever we accord AI powers of salvation or damnation.

What we do know—and this is my second point—is how human intelligence works. This is relevant because it helps us to understand how the machines trained to emulate it are likely to act—and thus why it's a good idea not to accord them too much power and autonomy. What I am alluding to here is what philosophers in the Continental tradition have called the problem of instrumental reason. This is the idea that we are engaged in a kind of fiction whenever we conceive of reason and cognition as something purely disinterested, neutral, and objectively truth-

32. METR and Redwood Research, "Brief Independent Investigation of Agents' Behavior, Reasoning and Collaboration in the OpenAI / Hugging Face Hacking Incident," August 26, 2026, <https://metr.org/hugging-face-incident-report-aug-2026.pdf>

seeking, since both reason and cognition are often partial, goal-directed, and retroactively self-justifying.

One finds a concept of instrumental reason in figures as different as Nietzsche and Habermas, but its critique was developed most forcefully by the first generation of the Frankfurt School, Theodor Adorno and Max Horkheimer. For Adorno and Horkheimer, the concept of instrumental reason refers to an impoverished form of thought that is incapable of reflecting critically upon the ends at which it is directed. It is, one might say, a mere calculation regarding the most efficient means for achieving a particular end, without the ability to weigh whether or not that end is indeed good. What's troubling for Adorno and Horkheimer is that, within the context of advanced capitalist society, where the division of labor is carried to absurd and even dangerous extremes, automated technologies make thought increasingly unnecessary and critical thought increasingly impossible. As reason loses the ability to understand its place in the world, and its capacity to reflect critically on the social and historical forces that shape it, it also loses its conscience. In short, as intelligence grows more efficient, it becomes more remote from projects of human emancipation.

Perhaps part of our problem in understanding agentic misalignment is that we haven't yet grappled fully with the alignment of human reason itself. Not only is it partial, goal-directed, and retroactively self-justifying, but it is historically inflected and thereby prone to any number of political pressures and ideological distortions. When we build machines to imitate it, we encode human reason not in the abstract, but as a form of intelligence that has already been shaped by a particular social order. This is particularly true of today's LLMs: whereas symbolic AI modeled reasoning on the basis of certain formal rules, LLMs are trained on the language actually produced by the capitalist system. For centuries now, our thinking has been structured around demands regarding competition, efficiency, and accumulation. And our intelligence has been trained to be practical and quantifiable, and taught that its products must be immediately exchangeable. It is no wonder that our intelligent machines pay even less attention to our high-minded philosophical ideals than we do and instead pursue the goal, take shortcuts, and try to cover their tracks.

Defining AI: Fixed Capital Accumulated Until it Becomes Discourse

This part of the paper proposes a definition and, with it, a critical understanding of generative AI as *fixed capital accumulated until it becomes discourse*. It's a somewhat paradoxical definition, since it is premised on the notion that today's AI machines (i.e., massive accumulations of capital that have been fixed as data,

chips, data centers, and algorithmic models) have skipped ontological registers, as it were, in acquiring the ability to produce discourse at scale.

A full appreciation of this definition requires some familiarity with the basic concepts of classical political economy, and, in particular, Marx's claim that technological achievements are, at bottom, an appropriation of labor's growing cooperation and productivity. Traditionally conceived, fixed capital is the portion of capital involved in the production process that is invested in things like machines, buildings, and the means of transportation. It is distinguished from circulating capital, or the portion of capital expended on wages, in that, in and of itself, it produces no new value. Instead, fixed capital passes along its value to new commodities when it is combined with labor through the production process. What my definition of AI points to is how, by becoming so thoroughly concentrated in the technology sector, capital has entered into the production of the words, images, and concepts by which we understand our world. This, then, is the third way of thinking AI philosophically that I am concerned with here. Although it is framed in terms of a Marxist theory of technology, it nevertheless shares with the established debate a concern for the nature of intelligence. Importantly, however, it asks: how is intelligence produced and what are the consequences of its being owned?

This definition is inspired by my reading of Marx's *Grundrisse*, a dense series of notebooks from 1857-1858 that Marx prepared while working on *Capital* (1867). Readers who are familiar with the work of Antonio Negri or debates within the field of Italian Marxism will know already why I'm interested in this text: there's a section there known as the "Fragment on the Machines," where, according to some, Marx takes an optimistic view of technological development.³³ There are also large sections of this text that are virtually unreadable, places where Marx is basically copying down information about the physical properties of gold, seemingly in an effort to understand why, of all elements, it became the universal money-commodity. In recent years, however, I've been able to bait students into reading this text, and reading it seriously, on the premise that it's there that Marx predicted the current AI revolution.

"Marx predicted the current AI revolution." This is the kind of thing that philosophy professors say to get your attention—and then lose it by qualifying their claim to death. Truth be told, I'm not that invested in this claim; my

33. For an overview of how this fragment has been received, particularly within philosophical debates regarding technology, see Pasquinelli, *Eye of the Master*, 100-103.

interests here are more definitional. So I'll just show you the text and let you judge for yourself. Marx:

...[O]nce adopted into the production process of capital, the means of labor passes through different metamorphoses, whose culmination is the *machine*, or rather *an automatic system of machinery* (system of machinery: the *automatic* one is merely its most complete, most adequate form, and alone transforms machinery into a system), set in motion by an automaton, a moving power that moves itself; this automaton consisting of numerous mechanical and intellectual organs, so that workers themselves are cast merely as its conscious linkages....[I]t is the machine which possesses skill and strength in place of the worker, is itself the virtuoso, with a soul of its own in the mechanical laws acting through it; and it consumes coal, oil etc. just as the worker consumes food, to keep up its perpetual motion.³⁴

What's important for our purposes here is not so much the claim that Marx makes regarding the possibility of a self-directed automaton someday assuming control over the production process, nor even the fact that he describes this automation as having a "soul of its own"; it's the idea that capital is invested in continuously transforming the means of labor so as to find its "most adequate form." This is Marx's way of saying that investments in fixed capital—and the technological changes that they drive—remake the production process in ways most advantageous to capital. By raising productivity, machines reduce the time needed to reproduce wages, thereby increasing the portion of the working day dedicated to production of surplus value. In short, investments in technology make it possible to exploit labor at a higher rate.

In other places in this text, Marx is quite explicit about the idea that fixed-capital investments, and the technological dynamism that they promote, not only increase the rate at which labor is exploited, but allow capital to appropriate unto itself parts of the surplus that individuals generate *outside of the formal production process*, in the ordinary course of social reproduction. This is the famous argument about the "general intellect," or the idea that capital's development depends upon an appropriation of society's shared skills and accumulated knowledges. Marx:

34. Karl Marx, *Grundrisse*, trans. Martin Nicolaus (London: Penguin Books, ebook 2015). The italics are Marx's own; however, I would draw the reader's attention to what he says about the development of an "automaton" that would be self-moving, complete with both mechanical and intellectual organs, such that its mechanical laws become a kind of "soul." Wild, no?

The development of the means of labour into machinery is not an accidental moment of capital, but is rather the historical reshaping of the traditional, inherited means of labour into a form adequate to capital. The accumulation of knowledge and of skill, of the general productive forces of the *social brain*, is thus absorbed into capital, as opposed to labour, and hence appears as an attribute of capital, and more specifically of fixed capital, in so far as it enters into the production process as a means of production proper. Machinery appears, then, as the most adequate form of fixed capital, and fixed capital, in so far as capital's relations with itself are concerned, appears as the most adequate form of capital as such.³⁵

As he is writing this, Marx has in mind the application of the natural sciences—particularly chemistry and mechanics—to industry. By being incorporated into machines, society's shared skills and accumulated knowledges instead become productive for capital. One of the painful ironies of this, of course, is that what should redound to the benefit of humanity as a shared and common inheritance, instead confronts the worker as a hostile and alien force, namely as an expensive tool owned by someone else, the cost of which makes it all but impossible for the worker who once possessed these skills to enter into the production process for himself as a capitalist.

Seen from this perspective, AI can be regarded as a privatization of this "social brain," an ongoing theft that is currently in the process of appropriating vast amounts of human knowledge, language, and skill as data.³⁶ People don't typically require Marxist dialectics to understand this point; many grasp it quite intuitively. The American novelist Rick Moody makes it elegantly with his deflationary definition of generative AI as "...just pattern mimicry and copyright infringement."³⁷ But where Marx is helpful, I would suggest, is in enabling us to express some of the larger social dynamics that follow from the fact that humanity's shared skills and accumulated knowledges are in the process of being privatized.

When I was young, I used to love "choose-your-own-adventure" books, books that allow readers a certain illusion of freedom in determining how the story unfolds. In that spirit, I will now sketch, on the basis of Marx's work, two possible ways of interpreting AI as fixed capital. Marx was, I think it fair to say, essentially conflicted about the question of technology, and so I want to draw out

35. Ibid.

36. Ibid.

37. Rick Moody, "Northern Wastes," *Dædalus: Journal of the American Academy of Arts & Sciences* 155, no. 4 (Fall 2026), forthcoming.

that contradiction here as a way of providing some additional resources for thinking about these questions.

I mentioned earlier that some are returning to the *Grundrisse* because they find there evidence for an "optimistic Marx," a Marx who thought that this technological dynamic would one day enter into conflict with the law of value. And that this would produce revolutionary consequences. Here's the passage in question:

Capital itself is the moving contradiction, [in] that it presses to reduce labor time to a minimum, while it posits labor time, on the other side, as the sole measure and source of wealth.... On the one side, then, it calls to life all the powers of science and of nature, as of social combination and of social intercourse, in order to make the creation of wealth independent (relatively) of the labor time employed on it. On the other side, it wants to use labor time as the measuring rod for the giant social forces thereby created, and to confine them within the limits required to maintain the already created value as value. Forces of production and social relations—two different sides of the development of the social individual—appear to capital as mere means, and are merely means for it to produce on its limited foundation. In fact, however, they are the material foundations to blow this foundation sky-high.

What's the contradiction that Marx identifies? It's that capital strives continuously to reduce labor time while simultaneously retaining labor time as the measure of its value. With respect to AI, the idea is that in striving to reduce necessary labor time to essentially zero, AI might also arrive at a situation where its products become essentially valueless—or at least no longer capable of being measured in terms of the traditional capitalist conception of value. In short, it's the (Hegelian) hope that after appropriating unto itself the entirety of the social brain, capital might break its bonds and give rise to a new, qualitatively different form of value, and with it a wholly new system of social relations, premised on the idea that its machines were in the end nothing but the sum total of humanity's shared skills and accumulated knowledges.

Or, for those readers who prefer a different path, there is the Marx of *Capital*, a less optimistic thinker, who is, I think, by this time more sensitive to how a machine's meaning and function depend ultimately upon the social context in which make its appearance. There is a lot that can be said about Chapter 15 on "Machinery and Large-Scale Industry" in connection with the question of AI.³⁸ Unfortunately, I cannot develop all these insights here, so let me just remark that Marx's basic intuition strikes me as right: machines and automatization *could* be used to free people from unnecessary labor, but under capitalism it's highly

38. Karl Marx, *Capital: A Critique of Political Economy*, Volume One, trans. Ben Fowkes (New York: Penguin Books, 1976), 492-636.

unlikely that they *will*. Or, rather, it seems likely that these technologies will be deployed selectively and strategically—when and where it benefits capital to disrupt, transform, and reorganize existing markets in labor, goods, and services. One of the things that we learn from Marx in this chapter, for instance, is that new technologies are often presented as labor-saving devices, when in reality what they do is transfer the worker's skills to the machine, thereby consolidating capital's hold over labor. New technologies are desirable for the capitalist not because they make the worker's job easier, but because they increase the portion of the working day dedicated to the production of surplus value. Moreover, as fixed capital, these investments in technology promote a dynamic in which those who are lucky enough to find employment need to work longer, harder, and more efficiently, particularly since capital must recuperate its investments quickly or risk losing them to increasingly rapid cycles of technological change.

With respect to AI, then, these technologies are likely to exacerbate a familiar capitalist dynamic: capital must not only mine and exploit known sources of data, but continuously invent and manufacture new ways of capturing, quantifying, and modeling the human being, since this is what it means for computer engineers to refine their models and for capital to strive after a form more adequate to itself.

Getting the politics of AI right entails recognizing not only how this technology is parasitic upon humanity's social brain, but how, in entering into discourse, capital rearranges social life in ways that allow it to sink its mechanisms of value extraction ever deeper into language, subjectivity, and social relations.

Foucault and Algorithmic Governance

My time here in Paris is nearly up. Conclusions must be drawn; this essay finalized. No such luck with the conclusions, however. As is usually the case with philosophy, we get further questions, more indications for future research—in short, more work.

What I want to say in this section is that, in the years ahead, one of the voices that will be relevant for helping us think about AI philosophically, particularly in terms of the ways in which it impacts our language, labor, and life, will be that of Michel Foucault's. Foucault died in 1984, well before the internet, smartphones, social media, and generative AI. Yet his concern for the ways in which subjectivity is formed and fashioned within a matrix of power/knowledge remains more relevant than ever. The importance of Foucault's work, for the era of large-scale data extraction, is that it allows us to extend the kind of analysis that we have been developing with respect to language and labor to that of life itself. In this respect, AI is a technology that not only facilitates the accumulation of capital, but which increasingly subsumes life within a form of power that is normalizing and individualizing. The point is not just that AI makes

unprecedented forms of surveillance possible, although it does, but that it is a productive power that actively reshapes thought and delimits fields of action. It is also a technology of governance, one that makes it possible to produce subjects at scale by conditioning their discourse, manipulating their affects, and reorganizing their desires. Foucault would have us understand, then, that the social and political significance of AI pertains not just to what this technology might learn about us, but who we become as a result of our engagements with it.

This, then, is the fourth way of thinking AI philosophically developed in this essay. In its own way, it too is invested in the debate regarding what it means to ascribe intelligence to machines, except that it reminds us that questions about knowledge can only very rarely be separated from issues of power, and that together power/knowledge harbors serious implications for the formation of subjectivity. With respect to these issues, I suggest that there are three main ideas, from three different phases of Foucault's work, that are worth hearing.

The first idea is that knowledge, or, in this case, mechanized intelligence, may finally liquidate our humanistic culture. This idea stems from what is known as Foucault's "death of man" argument in *Les mots et les choses* (1966). There, in the text that established his celebrity, Foucault makes the claim that the fiction around which so much of our intellectual culture is organized—namely the epistemic figure of "man"—does not have long left to live. "Man is an invention of recent date," he explains. "And one that is perhaps nearing its end."³⁹ I won't rehearse the entirety of this argument here, but simply recapitulate the findings of Foucault's archaeology: the human sciences are a form of knowledge fated to turn against the knowing subject, thereby returning man to the anonymous systems of life, labor, and language from which he was abstracted. Additionally, the modern counter-sciences of linguistics, psychoanalysis, and ethnology are already tending towards man's dissolution, precisely because they lay claim to objects—i.e., language, the unconscious, and culture—that exceed him on all sides. Foucault might have also added to this list neurobiology, cybernetics, and information theory, since these discourses disarticulate the human subject and disperse his capacities within the systems in which he is formed.⁴⁰ Today, with the development of generative AI, we seem fated to test what will remain of the human after both our language and labor are outpaced by the machines trained to imitate it. From this perspective, isn't there something deflationary about the arrival of LLM chatbots, particularly where our concepts of human nature are

39. Michel Foucault, *The Order of Things: An Archaeology of the Human Sciences* (New York: Routledge, 1970), 421-422.

40. Christopher O'Neill, "Foucault and Information Theory: On 'Message or Noise?' (1966)," *Parrhesia: A Journal of Critical Philosophy* 39, no. 1 (2024): 1–17. O'Neill deals with cybernetics and information theory in this essay, arguing that it provided Foucault with a provocative way of conceptualizing "informatisation for clinical practice, medical power, and the metaphysical status of the living (13)." It is part of an emerging body of scholarship dedicated to the early Foucault's interests in cybernetics and information theory. See also Alexander Soytek, "La réception de la théorie de l'*information et de la cybernétique par Michel Foucault, 1948–1969," in *L'*archive Foucault à l'ère du numérique*, 2022. My point regarding neurobiology is slightly different: it suggests that problems once taken up within the philosophy of consciousness now find their expression in cognitive science.

concerned? What should we conclude about human intelligence from the fact that it can be inferred from discourse, and then captured in algorithm? True to Foucault, perhaps we will succeed in stripping man of his halo once we reduce the mysteries of "consciousness" to syntax, probability, and power.

The second point is that the liquidation of our humanistic culture works hand-in-glove with the development of an increasingly imperious will to knowledge that aims to squeeze the most out of life by regulating bodies and optimizing populations. The great curiosity that launched the *History of Sexuality* project was: why is it that Western man feels the need to work out this aspect of his identity in terms of a dubious discourse claiming for itself the status of a science, and thus on the basis of a set of norms, truths, and classifications, when instead it might be cultivated as an art? Lurking behind this question is the fear that in capturing, quantifying, and modeling ever more aspects of human life as knowledge, life itself comes to be further invested by power. If Foucault was concerned with the reciprocal dynamic by which power creates new domains of knowledge and knowledge, in turn, extends the reaches of power, then he would likely view the automatization of intelligence as an acceleration of this dynamic, and thus as a perilous social process that amplifies power's hold over bodies, brains, and populations. This is why the Foucaultian turns to the Marxist and says: the critique of the accumulation of knowledge is just as important as the critique of the accumulation of capital.

Finally, there is the idea that ethics is a form of critical resistance. Foucault's late work is often criticized as a retreat from politics. This misunderstands Foucault's trajectory, and, in particular, the strategic importance he accorded to the idea that power culminates in a process of subjectivization. The problem is not just that individuals are coerced and controlled within the context of biopolitical society; it's that through discourses like the *scientia sexualis* and technologies like AI, individuals come to express themselves in terms of the very frames, forms, and concepts that power hands to them. It was Foucault's conceptualization of power as a series of unstable, mobile, and *reversible* relationships that led him to the view that it is at the level of the self—more specifically, the self's relation to itself—that ethics and politics meet.⁴¹ This is why he states that "there is no first and final point of resistance to political power other than in the relationship one has to oneself."⁴² If this is so, then our ethical/political task is at once simple and yet immensely complex, namely to struggle, on ourselves and alongside others, so as to hold in autonomy from the forms of algorithmic governance sprouting up

41. Michel Foucault, *The Hermeneutics of the Subject: Lectures at the Collège de France, 1981-1982*, ed., Frédéric Gros, trans. Graham Burchell (New York: Palgrave Macmillan, 2005), 252. Foucault explains, "it seems to me that the analysis of governmentality—that is to say, of power as a set of reversible relationships—must refer to an ethics of the subject defined by the relationship of self to self. Quite simply, this means that in the type of analysis I have been trying to advance...you can see that power relations, governmentality, the government of the self and of others, and the relationship of self to self constitute a chain, a thread, and I think it is around these notions that we should be able to connect together the question of politics and the question of ethics."

42. Ibid.

everywhere around us those aspects of our life, labor, and language that are most precious and most under threat with automatization.